

A nighttime photograph of a city skyline across a body of water. A prominent bridge with blue lighting spans the water in the middle ground. The city buildings in the background are illuminated with various lights, and their reflections are visible on the water's surface.

Application Architecture for the Cloud

Steve Loughran
Julio Guijarro

Slides:

<http://wiki.smartfrog.org/wiki/display/sf/Architectures+for+the+cloud>

© 2009 Hewlett-Packard Development Company, L.P.
The information contained herein is subject to change without notice





Why move to *the cloud*?

Good


- Cost
- Outsourcing hardware problems
- Move from capital to pay-as-you-go
- To handle Petabytes of data
- For a business plan that might work

Bad: to avoid the operations team

What is a cloud application?

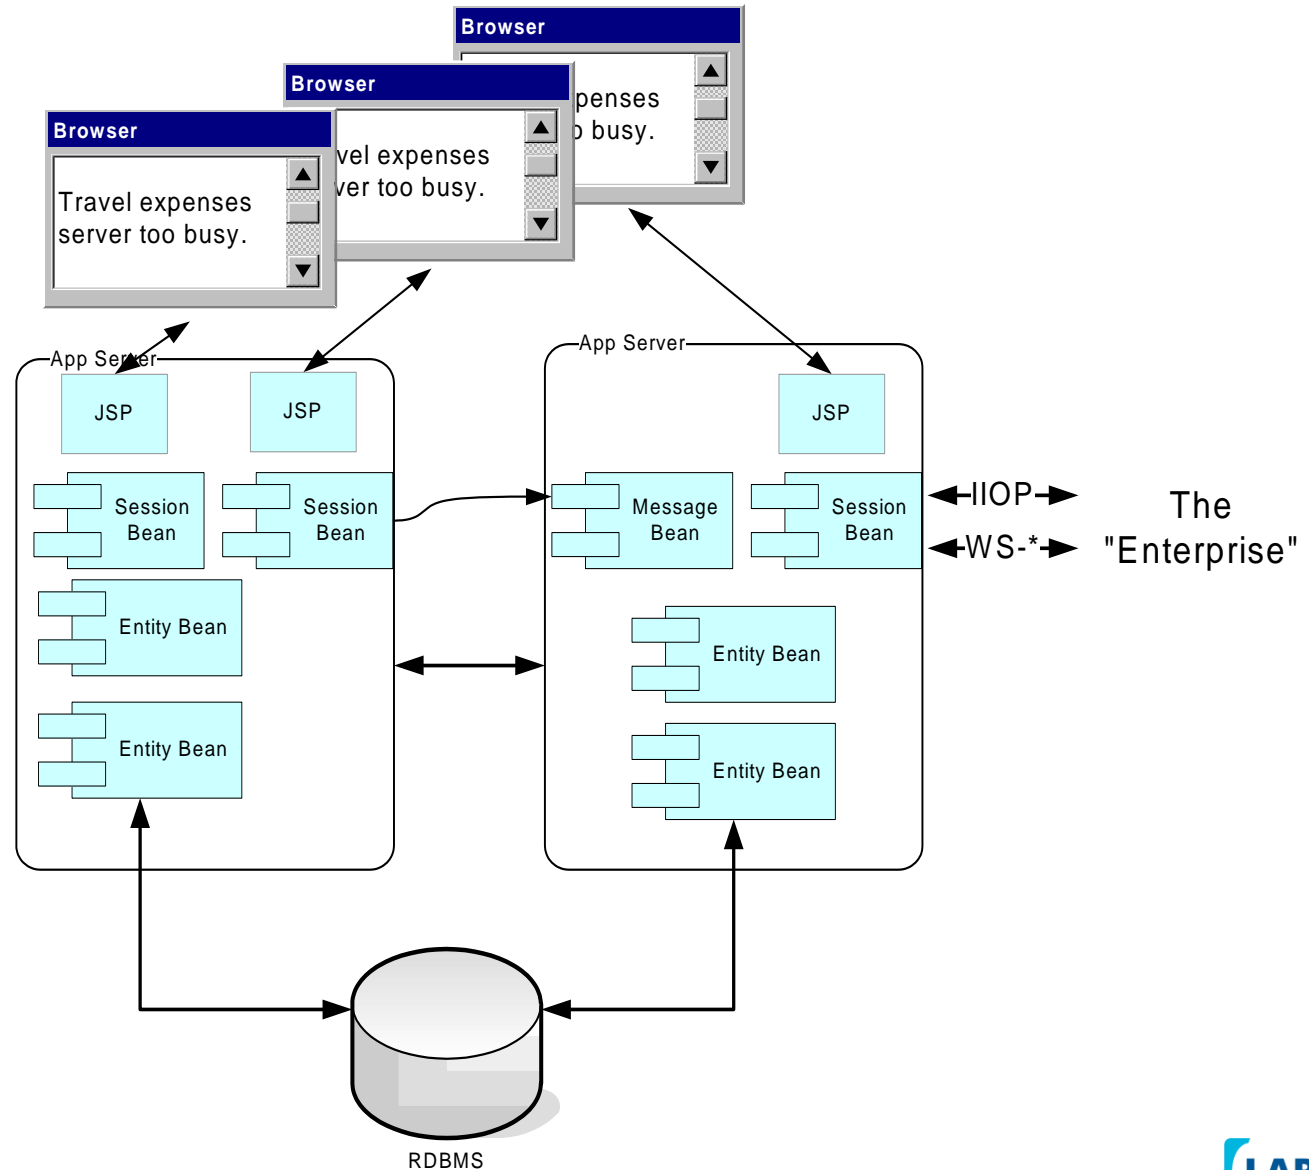
- The program that is run
- The code/data needed to run it in a datacentre
- Anything needed to configure, monitor and manage the system

This is not a cloud application



Enterprise Java on Highly Available servers

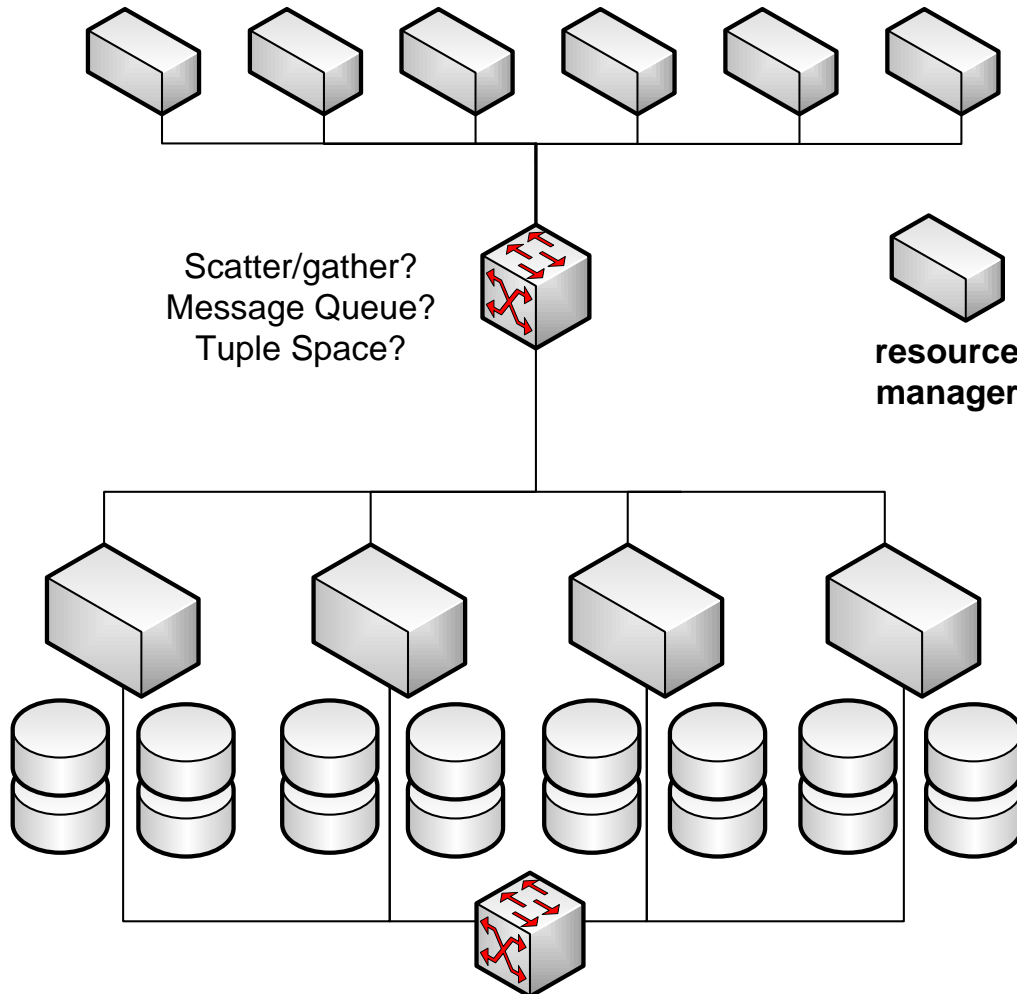
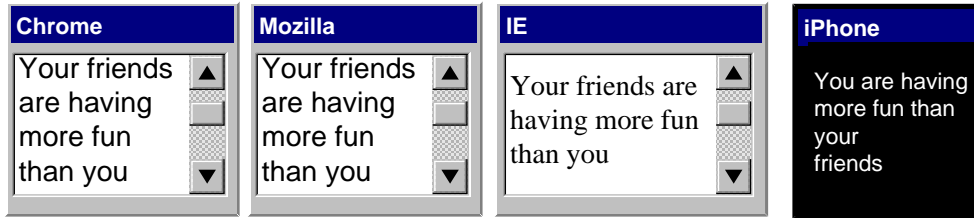
This is not a cloud application: Java EE



Things must change!

- Web UI for users, affiliates, marketing, operations
- Agile machine management is part of the API
- Scale up -and down
- Live upgrade of running system
- Persistence with key-value stores
- A Petabyte filesystem is part of the application
- MapReduce jobs close the loop
- Developers deploy to the cloud to test

REST APIs








Diskless front end

Memcached
JSP?
PHP?

Back End

HBase/Cassandra/
CouchDB
MapReduce
Lucene
HDFS or ...



EC2 as the cloud provider

- S3 for persistence, public downloads
- *SimpleDB* provides key-value storage, but query costs unpredictable.
- *Typica* for EC2 services API
<http://code.google.com/p/typica/>
- AWS IP rental or Dyndns for hostnames
- No image management services

No standard Apache Stack

Sun, IBM, HP as the cloud providers

- S3-like filestore
- EC2 and Sun RESTy APIs
- Unknown queue and keystore services
- More secure networking?
- Image management services?

No standard Apache Stack

Private cloud

- *Eucalyptus for deploying Xen images*
- Various persistence options
- Private filestore: HDFS, kfs, Lustre
- Kickstart for image management?

No standard Apache Stack

Whoever owns the API owns the application for its life

Whoever owns the data owns you

Apache Cloud Computing Edition

- Diverse mix of high-level technologies
- Very large filestore at the bottom :
Hadoop APIs, Java 7 NIO, Fuse, WebDAV
- MapReduce phase for post-processing
- We need stories for :
persistence, configuration, resource
management

A standard Apache Stack!

Front End

- The existing Web Front ends should work: Servlets, JSP, wicket, PSP, grails (maybe with memcached)
- Glue: queues, scatter-gather, tuple-space, events
- Everything needs to handle an agile world
- Everything needs instrumenting for management

Persistence



SimpleJPA



YAHOO!

PNUTS/Sherpa

Cassandra



Everything needs a REST API

- REST is the long-haul API -why have a separate internal one?
- Jersey JAX-RPC is very nice
- Client API evolving
- Http Components/HttpClient can be the foundation for the Apache client; needs AWS support.
- Also: *Restlet*

Events and messages

- "disk 3436 is failing"
- Bluetooth phone 04:5a:1f:c2:87:91 entered cell 56 in London NW2
- Queued purchases with card numbers

*internal and external events:
reliability, scalability, triggered actions*

Resource Management?

1. HA resource manager to monitor front end/back end load and request/release machines on demand
2. Kill unhealthy nodes (liveness, performance)
3. Programmable policies (money vs. load)
4. Choreograph live upgrade/migration
5. Resource Manager as a service

Configuration & Management

- LDAP (and APIs)
- key-value stores
- SmartFrog moving to Apache license
- What is Spring planning in this area?



Development

- How to build and test in this world?
- How to step through a program running on a remote datacentre?
- How to control testing costs?




Eclipse won the Java IDE battle - It now needs to compete with Visual Studio + Azure

Testing -your first terabyte of data

```
protected void map(Text key, Text test, Context context)
    throws IOException, InterruptedException {
    TestResult result = new TestResult();
    Class<?> testClass = loadClass(context, test);
    Test testSuite = JUnitMRUtils.extractTest(testClass);
    TestSuiteRun tsr = new TestSuiteRun();
    result.addListener(tsr);
    testSuite.run(result);
    for (SingleTestRun singleTestRun : tsr.getTests()) {
        context.write(new Text(singleTestRun.name),
            singleTestRun);
    }
}
```

Lots of opportunities here!

Testing -the infrastructure can help



Pseudo-RNG driven cluster configuration

Cirrus Cloud Testbed?

- HP, Intel, Yahoo!, universities
- Heterogeneous, multiple datacentres
- Offering datacentre time, not specific apps
- Low-level API for physical machines
- Cloud-API for virtual machines
- Paying customers? *No, not yet.*
- Open source projects? *I hope so*

What next?

Apache has the core of a Cloud Computing stack

How do we take this and:

- integrate the various pieces?
- extend them where appropriate?
- provide an alternative to AppEngine and Azureus?

Call to action

- Stop writing EJB apps
- Start collecting as much data as you can and feeding that MapReduce mining-phase.
- Design for: distributed not-quite-Posix filesystem, message queues, name-value databases

Apache: let's build our own cloud platform



VM Image Management

- AMI Image sprawl: 10%/month
- Old images are a security risk
- The whole PXE+Kickstart process is built for physical machines.

This is not an Apache problem, but we'll need to work with the OS vendors & others to integrate

HDFS improvements

- Scale, availability, small files: hierarchical namenodes?
- Could it be a general purpose media store?
- For web sites?